



AI Forge + AI Guild

ORG-WIDE AI ENABLEMENT · LIMS & SIGNATERA PLASMA

April 2026

Scaling AI adoption beyond a single *lighthouse team*.

Natera is ready to scale AI adoption beyond a single lighthouse team. Leadership across LIMS and Signatera Plasma wants to bring 20+ engineering teams to a consistent level of AI fluency — every PM, tech lead, engineering manager, developer, and SQE — and sustain that capability long after the engagement ends. The QCI Independence engagement proved what AI-native practices can do for one Natera team. This proposal is how Natera takes that playbook org-wide.

We start with a **2-team POC** over 3 weeks, pause for a joint retro with Natera engineers, then scale the refined playbook across the remaining 18+ teams. An **AI Forge** bootcamp each team cycles through, paired with an **AI Guild** — an active community of practice that ensures no one is dropped at the end of their bootcamp.

2 → 20+

2-team POC first, then scale to full org

~160

Engineers, SQEs, PMs, tech leads, EMs to enable

2-part

Engagement model: Forge + Guild

~3 mo

Guild Liatrio-led, then Natera-owned (Dave Peticolas)

THE ENGAGEMENT MODEL

The Forge and the Guild are *complementary by design*

Each is necessary, neither is sufficient on its own. The **Forge** is where new skills are forged under pressure: immersive, hands-on, shipping-real-work fluency built in 3 weeks. The **Guild** is where those skills keep growing after the cohort ends — a community of practice that sustains the muscle, spreads what's working across teams, and keeps an *innovative mindset* alive long after the bootcamp is over. Skills without a community decay; a community without new skills stagnates. Together they compound.

POC investment: \$289K — 2 teams, 3-week Forge, Guild stand-up in parallel.

AI Forge

3-week bootcamp · Per-team cohort · Ramp-up + 2 weeks immersion

A focused, hands-on AI enablement bootcamp each team cycles through. Rooted in Liatrio's internal **AI Forge Bootcamp** curriculum and calibrated for the full team: PM, tech lead, engineering manager, developers, and SQEs. Every role walks out using AI as a daily multiplier across their slice of the SDLC.

 **Who's in the cohort:** Developers · SQEs · Tech Leads · Engineering Managers · Product Managers

Curriculum · A Menu, Not a Commitment

The full flavor of possible focus areas across the SDLC. Each team's Forge is tailored from this menu based on a pre-Forge survey and team goals — not every topic is covered in every cohort.



Intake & refinement

AI-assisted PRDs, ticket shaping, Jira MCP workflows, Gantt roadmapping



Development

Claude Code setup, context management, parallel worktree workflows, spec-driven development



Code review

CI-integrated AI code review, deep review agents, HCLS-aware review patterns



Testing

BDD/Playwright, local E2E test suites, test generation, shift-left patterns for SQEs



Delivery flow

Branching, CI/CD observability, build-time optimization, merge workflow



Governance & safety

AI-generated code review expectations, prompt patterns, in-bounds/out-of-bounds for LIMS data

Core Three · Guaranteed in Every Cohort



Claude Code fundamentals (setup, context, worktrees) for all engineering roles



A team-specific Flashbuild outcome shipped against real backlog



At least one team-owned Claude skill or plugin living in the team's repo

3-Week Structure

WEEK 1 · RAMP-UP (25-50%)

Prerequisite course, async setup, and light collaboration while the team stays partly on their backlog. Handles setup friction *before* full immersion.

- Prereq course (Claude fundamentals)
- Environment & tooling setup
- Survey debrief & curriculum tailoring
- Initial pairing sessions


WEEKS 2 & 3 · IMMERSION (100%)

Full-time immersive AI enablement pairing. Learning happens on real backlog, culminating in a Flashbuild sprint.

- Role-specific labs (PM, Dev, SQE, TL, EM)
- Extended paired delivery on real tickets
- Full 3-4 day Flashbuild sprint
- Custom skill/plugin development
- Team-specific AI pattern library seeded

The Guild is the connective tissue between **Forge cohorts** and org-wide AI maturity.

A cross-team community of practice. The mechanism that makes sure no one is dropped after their bootcamp: daily office hours, open help blocks, a Natera-wide skills marketplace, and a Slack-native home for the people learning together.



Community ≠ Governance Body

Standards Are Emergent, Not Enforced

These are distinct functions. Combining them creates conflict. The Guild is primarily a community of practice. Governance is a separate, lighter-weight track the Guild *informs* — it does not own it.

The best tool or skill becomes the standard because it's the easiest and best — not because it's mandated. The Guild surfaces what's working across teams; adoption follows quality.

Aftercare, Made Concrete

- Daily office hours: on-call Champion rotation
- Open help blocks: scalable come-and-go sessions
- Natera-wide skills marketplace
- Slack-native: #ai-guild-wins, -help, -patterns

Early Initiative · OTEL Standardization

A concrete first win: standardize OTEL configuration across engineer laptops so AI usage telemetry (currently surfacing misleading numbers in Jellyfish) is actually accurate.

Liatrio spins up the Guild, develops 1-2 AI Champions per team via the Forge, and transitions full ownership to **Dave Peticolas** by the end of the engagement. The flywheel means the Guild gets stronger after handoff, not weaker.

FOUNDATION

The Guild's Three Pillars

Three reinforcing pillars built around community first. Standards emerge from what's working. Each pillar is Liatrio-seeded and Natera-owned.



PILLAR 01

Community & Cross-Team Learning

The Guild's primary identity. Every team's learning becomes every other team's advantage. Structured cadences, async channels, and aftercare that actually lands.

- Daily office hours: on-call Champion rotation
- Open help blocks: scalable come-and-go support
- Show & Tell: bi-weekly demos of real AI-assisted work
- Natera-wide skills marketplace

PILLAR 02

Emergent Standards & Shared Practice

Standards surface from what teams actually find useful — not from a top-down mandate. A living pattern library that becomes the standard because it's the easiest path.

- Living pattern library of skills, plugins, prompts
- Early OTEL initiative to fix Jellyfish signal
- Lightweight tool surfacing & review
- Informs (does not own) Natera AI policy

PILLAR 03

Innovation & Continuous Learning

Keeps the energy from the Forge alive. Prevents the Guild from becoming a static body.

- Forge alumni mentoring later cohorts
- Monthly challenge sprints (mini-hackathons)
- Advanced topics: agentic workflows, MCP, deep review
- External AI radar: monthly trend digest

OUTCOMES

What Each Team Walks Out With

Every team. Every role. Baseline fluency + a real shipped outcome — no one dropped at the end.

- ✔ Every role trained on practical AI workflows for their slice of the SDLC.
- ✔ Team-owned Claude skills and plugin configuration living in their repo.
- ✔ 1-2 AI Champions identified per team and seated in the Guild.
- ✔ A tangible Flashbuild outcome: real merged work shipped the new way.

ROLLOUT PLAN

POC First, Then Scale Across 20+ Teams

We start with a **2-team POC** over the 3-week Forge, then take a **1-week joint retro + iteration** with Natera engineers (not just Liatrio feedback) before scaling the refined playbook to the remaining 18+ teams in parallel waves. **We can move as quickly as Natera wants** — the number of teams running through the Forge in parallel scales to your appetite and capacity. Internal iteration loops are how we'd operate regardless of scope. Exact sequencing and team selection finalized with Greg, Aimee, Dave Peticolas, and the LIMS/Signatera VP sponsors during kickoff.

POC	2 teams	Month 1 (3 weeks) · First Champions identified · Guild stands up in parallel
Retro + Iterate	1 week	Week 4 · Natera engineer feedback · Refine playbook & curriculum menu
Waves 2–4	Scales to Natera's pace	Months 2–4 · Parallel cohort waves sized to Natera's appetite · Guild transitions to Natera ownership
Waves 5–7	Scales to Natera's pace	Months 4–6 · Natera-owned Guild (Dave Peticolas) supports remaining waves · Liatrio in advisory mode

Directional. We can move as quickly as Natera wants — the number of teams running through the Forge in parallel scales to your appetite, capacity, and preferred rollout pace. Wave size, POC team selection, and sequencing tuned to team readiness and shared dependencies.

PROOF POINT

QCI Independence · 8-Week Engagement

These are not projections. These are results. The QCI Independence engagement was an 8-week embedded partnership with a Natera engineering team. The Forge + Guild model is designed to reproduce these outcomes across 20+ teams and make them sustaining.

Headline Results



Value Stream Transformation

End-to-end delivery cycle improvements, not just individual developer workflows.



Flashbuild Sprint · 4 Days Onsite

A 4-day focused sprint demonstrated what the team could do once the new practices were in place. Git worktrees with parallel Claude Code sessions split work into concurrent streams. 8× throughput increase — real product work, not exercises. Each 3-week Forge includes a Flashbuild sprint of this shape during Weeks 2–3 immersion.



In Their Own Words



"The Liatrio team has been instrumental in helping us leverage AI tools to improve overall team efficiency... they have also helped shift our approach from primarily manual development toward a more automation-driven mindset."

Jignesh Parekh
Senior Engineer, QCI



"I have been introduced to new ways of doing things that I might not have otherwise dove into... no question has been received negatively, which has only empowered me to try more and more."

Disha Chaubey
Technical Product Manager, QCI



"We learned many valuable new skills on workflows and plugins... leveraging these methods in our team, and sharing the gained knowledge with other teams has the potential of greatly increasing efficiencies across the organization."

Marc Fasnacht



"It was really an eye opening experience... I learnt a lot of sleek solutioning using Claude skills. These certainly have helped me direct Claude better, utilize agents, worktrees for multitasking and reimagine testing workflows."

Anu Jayaraman

BUSINESS CASE

Eight-Figure Annual Value at Steady State

📈 ROI Logic (QCI-Calibrated)

The QCI team saw a **49% reduction in total cycle time** on a fully loaded team spend of roughly **\$1.25M/year**. Applying a conservative **floor-case 30% effective productivity gain** per team across 20 teams puts annual value recovered well into the **eight-figure range** at steady state. Capacity creation — new products shipped without new headcount — compounds on top. Governance value (reduced AI risk in a regulated environment) is harder to quantify but meaningful.

\$1.25M

Annual team labor spend (per team)

30%

Floor-case productivity gain

20+

Teams enabled

8-figure

Annual value at steady state

Specific scenarios and targets firm up during kickoff discovery.

🛡️ The Engagement Philosophy

AI-Native Across the SDLC

Every role, every stage. PMs shape requirements with AI. Devs pair with Claude. SQEs build test suites and review coverage with agents. EMs use AI for reporting and unblocking.

Learn by Shipping

Every Forge cohort produces a tangible Flashbuild outcome on real backlog. No sidebar workshops. The learning is the work.

Governance From Day One

In a healthcare context, AI standards can't be an afterthought. The Guild codifies safe patterns, approved tools, and review expectations before drift accumulates.

Liatrio Makes It Stick

Liatrio spins up, upskills Champions, runs the Guild, then steps back. The flywheel design means the Guild gets stronger after handoff, not weaker.

Cost of Inaction

Without a Forge + Guild model, the most predictable outcome isn't "no AI" — it's **fragmented AI**. 20 teams, 20 different AI patterns, 20 different approaches to handling sensitive data. In a healthcare environment, that is a **compliance risk** before it is a productivity problem.

And the productivity gap between AI-fluent and AI-lagging teams will only widen. Every month of delay is **capacity left on the table** and risk allowed to accumulate.

PREREQUISITES

Prerequisites for Success



Executive Sponsor & Air Cover

A named VP-level Executive Sponsor (expected: VP of Engineering, LIMS +



Protected Team Time

Teams in their Forge window commit ~25-50% in Week 1 (prereq + ramp-



Environment & Tooling Access

Uninterrupted access to application environments, repositories, GitLab

Signatera Plasma alignment) who protects team time for the Forge, reviews Guild health quarterly, and removes blockers.

up) and ~100% in Weeks 2 & 3 (immersion). Guild Champions commit 2-4 hrs/week ongoing. Wave sequencing planned around release cadences.

CI/CD, Jira, Confluence, and existing telemetry platforms. Claude Code + Bedrock access confirmed (as used during QCI).

CHANGE MANAGEMENT

A Path, Not a Flip

Not every engineer, SQE, or PM arrives at their Forge fully sold on AI-native ways of working — that's expected. The Forge is designed as a structured opt-in: paired work, role-specific labs, a Flashbuild outcome people can point to, and a Champion path for those who want to lead. The Guild catches people post-Forge with cadences that are **social, not bureaucratic**. Leadership aligns up front on what success looks like for each role, and the proof points come from the cohort's own shipped work. The goal is a community that *genuinely owns the new way*, not compliance.

RECOMMENDED NEXT STEPS

From Proposal to Kickoff

01 Align on Approach

Confirm Forge + Guild as the engagement model. Lock 3-week Forge structure (Week 1 ramp-up at 25-50%, Weeks 2-3 immersion at 100%).

02 Select POC Teams

Pick the 2 POC teams across LIMS and Signatera Plasma. Name the Executive Sponsor and confirm Dave Peticolas as the Guild ownership recipient.

03 Run Pre-Forge Survey

Survey POC teams on AI usage, pain points, and goals so Liatrio can tailor the curriculum menu.

04 POC + Retro

Run the 3-week Forge with 2 teams; Guild stands up in parallel. Pause for a joint retro with Natera engineers.

05 SOW for Scale-Out

Refined SOW with deliverables, wave schedule, Guild operating model, and pricing for the remaining 18+ teams.

06 Scale

Waves roll out in parallel; Dave Peticolas owns the Guild by engagement close.



Confidential. Not for reproduction or further distribution.